

Topic 1: Introduction to Tensorflow


Tensorflow is a Python library for performing symbolic mathematical operations, which means that you create a graph structure that represents a number of connected mathematical operations that can be manipulated to do things like differentiation. It can also automatically evaluate the graph structure on the GPU for you, which makes it faster.

It is important to know what a tensor is: a structure of numbers organised as a grid with a certain number of dimensions. A 0-dimensional tensor is called a scalar, a 1-dimensional tensor is called a vector, a 2-dimensional tensor is called a matrix, a 3-dimensional tensor is called a 3tensor.

1. Graph and session

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/01 - Graph and session.py>

Tensorflow works using a graph and a session to evaluate the graph. Sessions take in inputs to the graph and return outputs. Here is the Tensorflow graph with Tensorflow nodes as rectangles and Python variables as ellipses.



Every time you use a node in an operation, a new node is created and added to the graph. You can visualise this graph in TensorBoard by running `'tensorboard --logdir=<path to folder with python file>'` in the terminal and then opening the URL shown in the terminal (do this after running the program).

2. Tensor expressions

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/02 - Tensor expressions.py>

You can compute expressions that work on tensors of any dimension instead of scalars. Note that the equation in the code is using an explicit conversion of Python constants into Tensorflow constants. We could instead let it happen implicitly (automatically) by writing `'2*a + 1'` but there would still be a conversion.

3. Getting/setting intermediate nodes

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/03 - Getting-setting intermediate nodes.py>

You can set a value for any intermediate node, as well as get the value of any intermediate node.

4. Variably sized tensors

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/04 - Variably sized tensors.py>

Tensors do not need to have a fixed size. Setting a size to 'None' will allow you to use any number of elements in that dimension. Any other expressions that depend on that tensor will change their size accordingly.

5. Variables

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/05 - Variables.py>

It is usually the case in machine learning that we need to use a model with parameters that are adjusted and learned so that we get a desired behaviour, that is, how inputs are mapped into outputs. In Tensorflow these parameters are called variables. Variables are not passed in with the inputs but **are stored within the session**. They are not the variables in the sense of Python's variables but are special nodes in the graph. There are also special nodes for setting variables and the values set for the variables are stored in the session.

6. Initialisers

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/06 - Initialisers.py>

Multiple variables can be initialised together using an initialiser node.

7. Variable names and scopes

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/07 - Variable names and scopes.py>

Variables must have unique names in order to be able to refer to them. In order to avoid coming up with long names we can organize them into variable scopes which act like folders. A variable in one scope can have the same name as another variable in another scope. It is possible to get a variable by its name by using 'get_tensor_by_name'. Every variable name ends with a ':0' such as 'varname:0'. If the variable is in a scope then its name will be 'scopename/varname:0'. If the scope is nested within another scope then its name will be 'scopename1/scopename2/varname:0'.

8. Saving and restoring variables

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/08 - Saving and restoring variables.py>

You can save and load variables for later use by using a 'saver' graph node. Remember that variable values are session specific, so closing a session will destroy the variable values. If you want to use the same variable values in a different session you will need to either manually set them or load them from a saved model. The saver node is useful for publishing or loading a trained model or even for saving your half-trained model's variables in order to resume training if the program crashes or if the power goes out.

9. Polynomials

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/09 - Polynomials.py>

Here is an example of a model consisting of a polynomial.

10. Plotting with matplotlib

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/10 - Plotting with matplotlib.py>

Plotting a polynomial with matplotlib.

11. Gradients

<https://github.com/mtanti/deeplearningtutorial/blob/master/tensorflow/01 - Introduction to Tensorflow/11 - Gradients.py>

Gradients are very useful in machine learning as they help you find the lowest point in a function or even the most sensitive input in a function (a high gradient means that a small change in the input will give a big change in the output). Tensorflow has functions for automatically finding the gradient of a graph. Unfortunately, we can only find the gradient of a scalar. It can be with respect to any shape tensor but the node we are finding the gradient of must be a single number, otherwise the gradient function will find the gradient of the sum of all elements in the tensor (it does not produce a Jaccard matrix).